


# autobox

A small computer that does software work on its own

It's a used mini desktop with no screen and no keyboard, and AI agents run on it all day. Its owner messages it from a phone, the way you'd message a colleague, and gets finished work back.

One person set it up. It now builds and runs several projects, among them a tutoring site and a circuit board designer, and most days it works on itself too.

FIGURE 1 *one message, and the work that comes back*




Slate is the person, ochre the agent that decides, sage the agents that do the work, and rust the checks every change passes. The dashed line is the note that goes back when it is done.

## Why it can be left alone

An agent that writes code is easy to get now.

No agent merges its own work. Every change goes through the box's test suite, and then a second agent reads it against what was asked for. If either one objects, the change goes back to the agent that wrote it, with the reason.

FIGURE 2 *what a change goes through before it merges*



Sage is the worker, rust the two checks, slate the merged result. The dashed line is the way back.

The owner is only asked about a short list of things, like spending over budget or anything that opens the machine up to the outside. The agents decide everything else for themselves.

Other people use it too. Each of them gets a private chat channel and a workspace of their own, and their agent can't reach anyone else's files.

When something goes wrong, it gets written down, and a mistake that keeps coming back turns into a new check.

Figure 2 is the path a change takes through the landing log at `~/cc/state/land/queue.log`: of the 443 changes that reached the merge queue between 1 and 21 September 2026, 152 were sent back at least once.

---

## Its first four weeks

These numbers come from the box's own records, read on 21 September 2026. The first line of its code was written on 26 August.

---

741

changes merged across eight projects in 26 days. That's about 28 a day.

1 in 3

changes were stopped at least once by the box's own checks before they merged. That's 152 of the 443 that reached the merge queue since 1 September.

97

of those 152 were stopped by the reviewing agent after the tests had passed.

8

people have their own workspace on it.

---

65 W

is the size of its power supply. It's a six core office desktop with 32 GB of memory, and the models run at the AI provider, so the box itself can be small.

62

small tools make up the system, about 65,000 lines. The box's own agents look after them, through the same checks.

---

The generic part of the system is published as open source under the name autobox. This box's own configuration stays private.

## Where this sits

Agents can already work for hours. The open question is whether you can trust what comes back.

*What the field has.* METR measures how long a task an AI can finish by itself. In January 2026 its best measured model managed tasks that take a skilled person about five hours, and that length has been doubling roughly every three months.<sup>1</sup> Anthropic, OpenAI, Google, GitHub and Cursor all sell agents that work in the background and hand back a pull request.<sup>2</sup> The top score on the standard coding benchmark is about 97%.<sup>3</sup>

*What it does not have.* The same leaderboard warns that top scores may be inflated, because the test has leaked into training data.<sup>3</sup> In a randomized trial by METR, experienced developers using AI tools took 19% longer on real tasks, and still believed they had been faster.<sup>4</sup> Researchers note that several agents working together fail in ways a single one does not, and that current safety frameworks do not cover it.<sup>5</sup>

---

## The hard part is being able to leave it alone while it works, and most of this project is what makes that safe.

THE ONE CLAIM THIS DOCUMENT MAKES

---

WHAT THIS DOES NOT  
MEAN

A change that passed is not a change worth making: the test suite and the reviewing agent check it against what was asked for, not whether the asking was right.

---

*Where the box sits.* The box doesn't try to build a smarter agent. It uses the ones the labs sell, and it can move to a better one when that arrives. What it adds is the layer around them, deciding who may do what, how work is checked, what it costs and what happens when an agent is wrong. The products above sell you the agent. This box runs many of them across several projects, on hardware the owner controls, and other people use it too. It has done that every day for nearly four weeks, and that's the record it has to show.

FIGURE 3 *the three layers, and which one this project is*



Slate is the people at the top and sage the agents this box buys at the bottom. The rust layer between them is the box.

---

## Sources

All sources opened on 21 September 2026.

- 1 METR, Time Horizon 1.1, 29 January 2026. [metr.org/blog/2026-1-29-time-horizon-1-1](https://metr.org/blog/2026-1-29-time-horizon-1-1)
- 2 Product pages: [claude.com/product/claude-code](https://claude.com/product/claude-code), [learn.chatgpt.com/docs/cloud](https://learn.chatgpt.com/docs/cloud), [blog.google/technology/google-labs/jules](https://blog.google/technology/google-labs/jules), [github.blog/changelog/2025-05-19-github-copilot-coding-agent-in-public-preview](https://github.blog/changelog/2025-05-19-github-copilot-coding-agent-in-public-preview), [cursor.com/help/ai-features/background-agents](https://cursor.com/help/ai-features/background-agents)
- 3 SWE-bench Verified leaderboard, updated 4 September 2026. [leaderboard.steel.dev/leaderboards/swe-bench-verified](https://leaderboard.steel.dev/leaderboards/swe-bench-verified)
- 4 METR, study of experienced open source developers, 10 July 2025. 16 developers, 246 tasks. [metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study](https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study)
- 5 arXiv 2603.09002, March 2026.